




Release Engineering for ML Applications




Introduction

1. ML Training Pipeline
2. ML Testing Design
3. Release Pipeline
4. K8s Deployment
5. Experimental Setup: Istio & Prometheus
6. Additional Istio Use Case
7. Extension Proposal

1. ML Pipeline

- DVC
- AWS S3
- Poetry package/dependency manager



1. ML Pipeline: artifacts

- Trained model
- Raw and processed data
- metrics

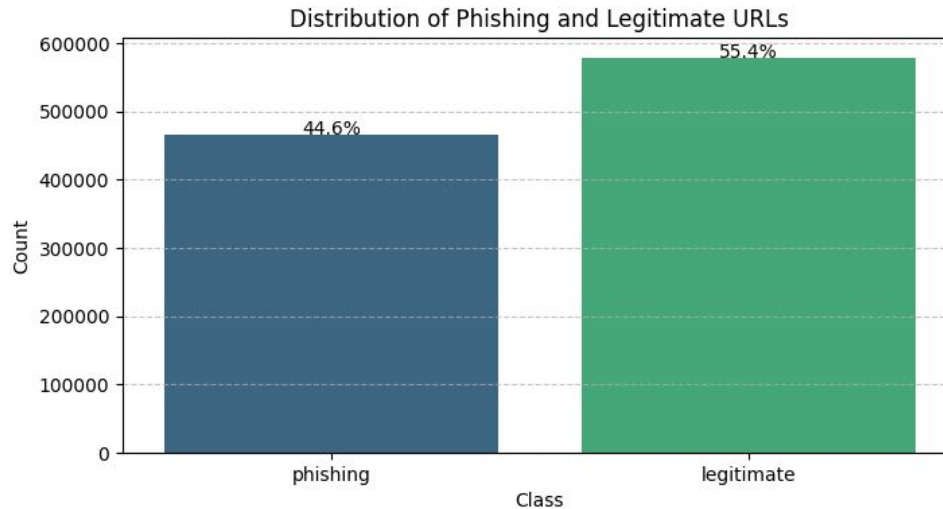
Metric	Value (%)
Accuracy	0.00022
Val Accuracy	0.000002
Precision	44.6082
Val Precision	44.6080
Recall	99.9593
Val Recall	100.0000
Loss	0.35034
Val Loss	0.35619

1. ML Pipeline: linters

- Pylint & Flake8
- Allow common naming conventions
- Discourage 'bad' naming
- Handle more exceptions
- Display only warnings with high confidence levels


2. ML Testing Design: features and data

- View distribution of the data



2. ML Testing Design: Features and Data

- View distribution of the data
- Engineered features

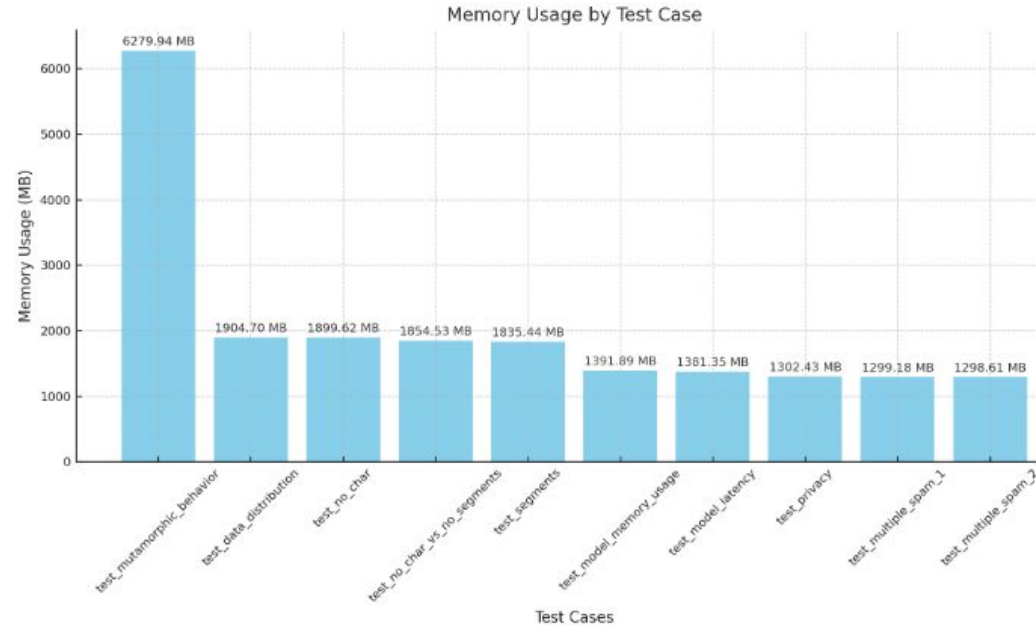


2. ML Testing Design: Mutamorphic tests

- Check for non-deterministic behaviour

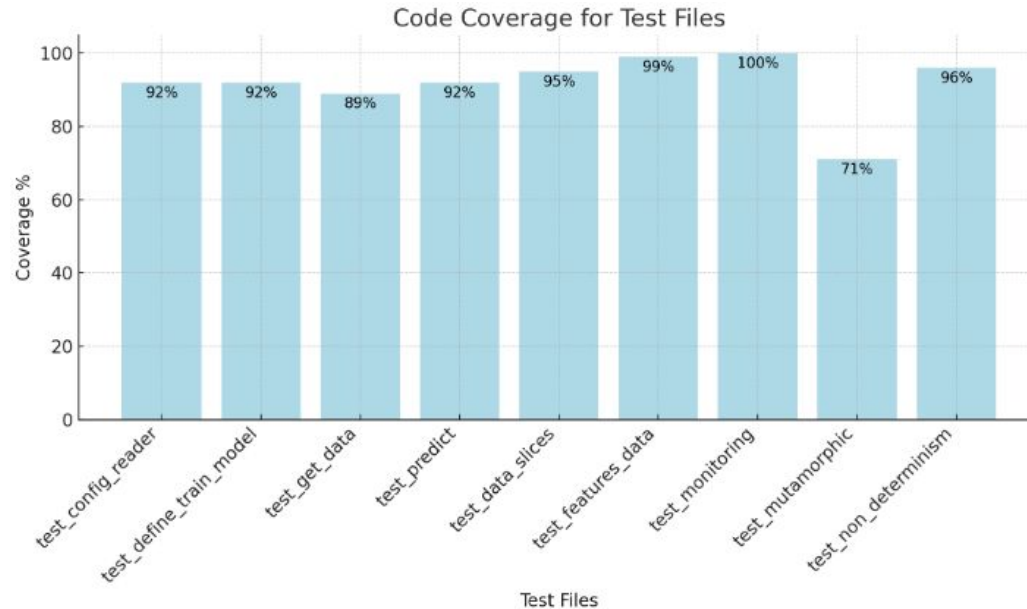
2. ML Testing Design: Monitoring Tests

- Ram usage per method

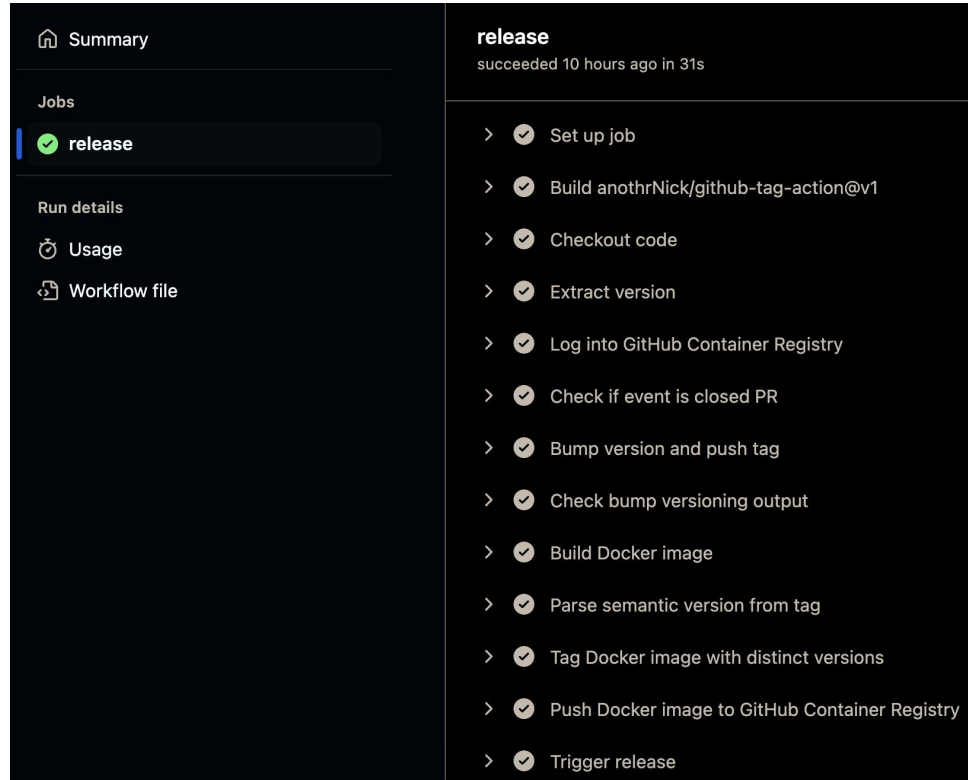


2. ML Testing Design: Test Adequacy

- Statements, covered, missed



3. Release Pipeline: Container Image (app)




The screenshot displays the GitHub Actions interface for a workflow named 'release'. The left sidebar shows the 'Summary' tab selected, with a list of jobs including 'release' which is marked as successful with a green checkmark. Below the jobs list, there are links for 'Run details', 'Usage', and 'Workflow file'. The main panel shows the 'release' job details, indicating it 'succeeded 10 hours ago in 31s'. A list of 14 steps follows, each with a green checkmark icon and a right-pointing chevron, indicating successful completion. The steps include setting up the job, building a Docker image, pushing it to GitHub Container Registry, and triggering the release.

release
succeeded 10 hours ago in 31s

- > ✓ Set up job
- > ✓ Build anothrNick/github-tag-action@v1
- > ✓ Checkout code
- > ✓ Extract version
- > ✓ Log into GitHub Container Registry
- > ✓ Check if event is closed PR
- > ✓ Bump version and push tag
- > ✓ Check bump versioning output
- > ✓ Build Docker image
- > ✓ Parse semantic version from tag
- > ✓ Tag Docker image with distinct versions
- > ✓ Push Docker image to GitHub Container Registry
- > ✓ Trigger release

3. Release Pipeline: Container Image (app)

```
app / .github / workflows / release.yml   
  
nadinekuo Update CI  
  
Code Blame 108 lines (94 loc) · 4.45 KB  
  
1  name: Release to GitHub Container Registry  
2  
3  on:  
4    pull_request:  
5      branches:  
6        - main  
7      types:  
8        - closed  
9    push:  
10     tags:  
11       - 'v[0-9]+.[0-9]+.[0-9]+' # Semantic versioning (e.g., v1.0.0, v2.3.4)  
12  
13  jobs:  
14    version-and-release:  
15      runs-on: ubuntu-latest  
16      permissions:  
17        contents: write # Needed for automatic patch bumping  
18
```

3. Release Pipeline: Container Image (app)

```
18
19     steps:
20       - name: Checkout code
21         if: github.event.pull_request.merged == true
22         uses: actions/checkout@v3
23         with:
24           ref: ${ github.event.pull_request.merge_commit_sha }} # Needed for automatic pre-releases
25           fetch-depth: '0'
26
27       - name: Checkout code
28         if: github.event.pull_request.merged != true
29         uses: actions/checkout@v3
30
31       - name: Extract version
32         id: version          # Set ID to refer back to in later steps
33         run: echo "::set-output name=version::${GITHUB_REF#refs/tags/}"
```

3. Release Pipeline: Container Image (app)

```
35     - name: Log into GitHub Container Registry
36       run: echo "${{ secrets.TOKEN }}" | docker login ghcr.io -u "${{ github.actor }}" --password-stdin
37
38     # In the case of closing PRs, automatically bump current version (default: patch)
39     # Any commit message that includes #major, #minor, #patch, or #none will trigger the respective version bump.
40     # If no #major, #minor or #patch tag is contained in the merge commit message, it will bump whichever DEFAULT_BUMP is set to
41     - name: Bump version and push tag
42       if: github.event.pull_request.merged == true
43       id: bump
44       uses: anothrNick/github-tag-action@v1
45       env:
46         GITHUB_TOKEN: "${{ secrets.GITHUB_TOKEN }}"
47         WITH_V: true           # Tag with v character
48         DEFAULT_BUMP: patch   # Which type of bump to use when none explicitly provided in commit msg
49         PRERELEASE: true      # Define if workflow runs in prerelease mode
50         PRERELEASE_SUFFIX: beta # Suffix for your prerelease versions. Note this will only be used if a prerelease branch.
```

3. Release Pipeline: Container Image (app)

```
60     - name: Build Docker image
61       run: docker build -t app .
62
63     - name: Parse semantic version from tag
64       id: semantic-version
65       run: |
66         echo "GITHUB_REF: ${GITHUB_REF}"
67         if [[ "${{ github.event.pull_request.merged }}" == 'true' ]]; then
68           echo "Case PR closed: the automatically bumped version will be used"
69           VERSION=${{ steps.bump.outputs.new_tag }}
70         else
71           echo "Case pushed tag: the manually created tag will be used"
72           VERSION=${GITHUB_REF:11}
73         fi
74         echo "$VERSION"
75         MAJOR=`echo "$VERSION" | cut -d . -f 1`
76         MINOR=`echo "$VERSION" | cut -d . -f 2`
77         PATCH=`echo "$VERSION" | cut -d . -f 3`
78         echo "::set-output name=version::$VERSION"
79         echo "::set-output name=version_major::$MAJOR"
80         echo "::set-output name=version_minor::$MINOR"
81         echo "::set-output name=version_patch::$PATCH"
82
```

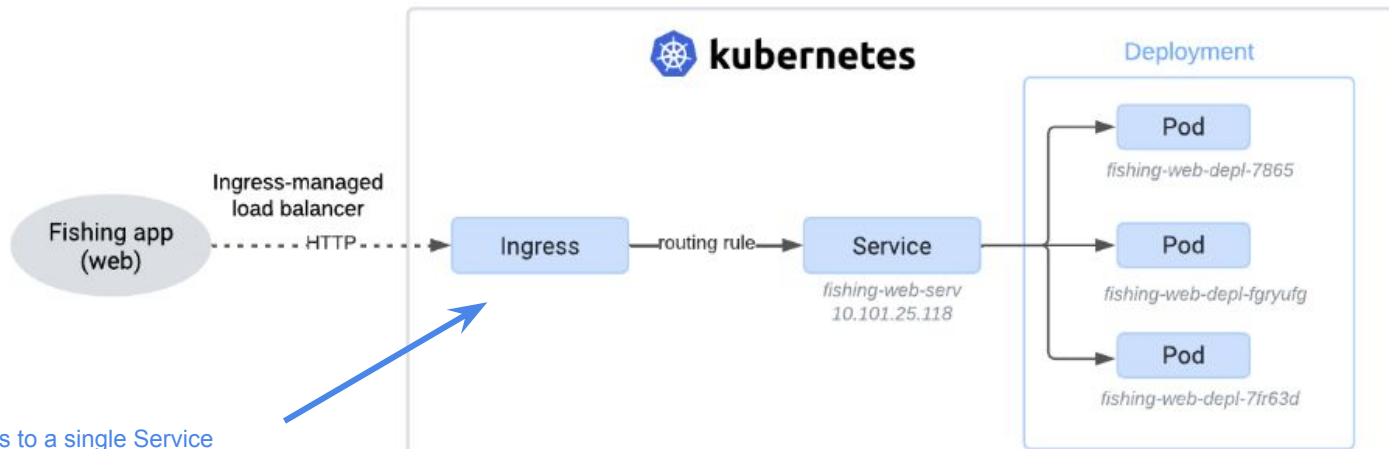
3. Release Pipeline: Container Image (app)

```
83   - name: Tag Docker image with distinct versions
84     run: |
85       IMG=ghcr.io/${{ github.repository }}
86       docker tag app $IMG:${{ steps.semantic-version.outputs.version }}
87       docker tag app $IMG:latest
88       docker tag app $IMG:${{ steps.semantic-version.outputs.version_major }}.${{ steps.semantic-version.outputs.version_minor }}.latest
89       docker tag app $IMG:${{ steps.semantic-version.outputs.version_major }}.latest
90
91   - name: Push Docker image to GitHub Container Registry
92     run: |
93       IMG=ghcr.io/${{ github.repository }}
94       docker push $IMG:${{ steps.semantic-version.outputs.version }}
95       docker push $IMG:latest
96       docker push $IMG:${{ steps.semantic-version.outputs.version_major }}.${{ steps.semantic-version.outputs.version_minor }}.latest
97       docker push $IMG:${{ steps.semantic-version.outputs.version_major }}.latest
98
99   - name: Trigger release
100     uses: actions/create-release@v1
101     env:
102       GITHUB_TOKEN: ${{ secrets.TOKEN }}
103     with:
104       tag_name: ${{ steps.version.outputs.version }}
105       release_name: Release ${{ steps.version.outputs.version }}
106       body: |
107         Release app version ${{ steps.version.outputs.version }}
```


3. Release Pipeline: Software Package (lib-ml)

```
53     - name: Set up Python
54       uses: actions/setup-python@v2
55       with:
56         python-version: '3.8'
57
58     - name: Install dependencies
59       run: |
60         python -m pip install --upgrade pip
61         pip install setuptools wheel twine
62
63     # Both source distribution (.tar.gz or .zip) and built distribution (.whl) are built
64     - name: Build Distribution
65       run: |
66         python setup.py sdist bdist_wheel
67
68     # Upload all files to PyPi
69     - name: Publish lib to PyPI
70       if: startsWith(github.ref, 'refs/tags/v')
71       run: |
72         twine upload dist/*
73     env:
74       TWINE_USERNAME: __token__
75       TWINE_PASSWORD: ${ secrets.remla_a2 }
```

4. K8s Deployment

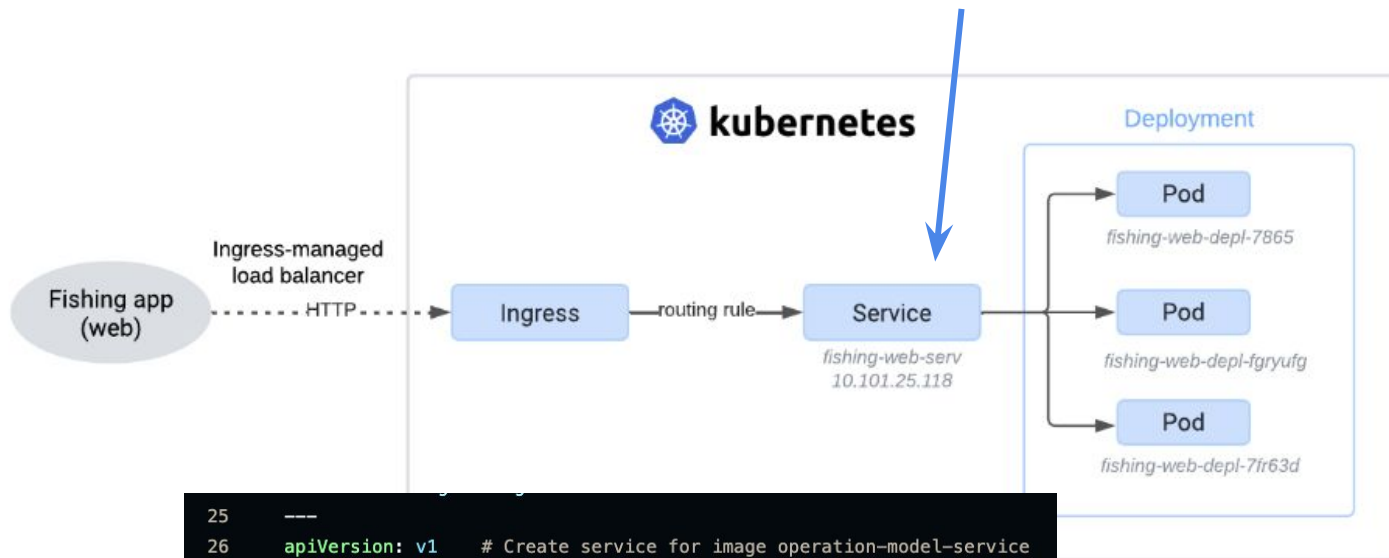


- Routes requests to a single Service
- Ingress Controller determines relevant Service port

```
79 ---
80 apiVersion: networking.k8s.io/v1 # Create an Ingress for the application (to expose the ClusterIP service)
81 kind: Ingress
82 metadata:
83   name: fishing-gateway
84 spec:
85   defaultBackend:
86     service:
87       name: fishing-web-serv
88       port:
89         number: 8000
```

4. K8s Deployment

- Prevent internal Pods from being accessed publicly
- Listen on specified port (fixed)
- Forward requests to containers in Pods

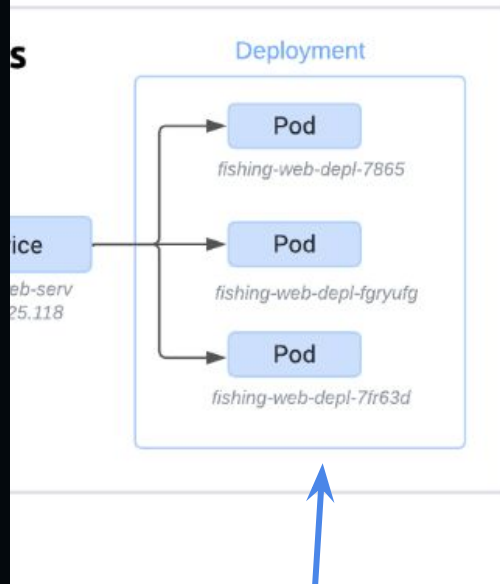


```
25 ---
26 apiVersion: v1 # Create service for image operation-model-service
27 kind: Service
28 metadata:
29   name: fishing-model-serv
30 spec:
31   selector:
32     app: fishing-model # App defined in Deployment above
33   ports:
34     - port: 5001
35       targetPort: 5001
```

4. K8s Deployment

Fishing app
(web)

```
36 ---
37 apiVersion: apps/v1      # Create deployment for image operation-app
38 kind: Deployment
39 metadata:
40   name: fishing-web-depl
41   labels:
42     app: fishing-web
43 spec:
44   replicas: 1
45   selector:
46     matchLabels:
47       app: fishing-web
48   template:
49     metadata:
50       labels:
51         app: fishing-web
52     spec:
53       containers:
54       - name: fishing-web
55         image: ghcr.io/remla2024-team14/app:latest
56         imagePullPolicy: Always
57         ports:
58         - containerPort: 8000
59         env:
60         - name: MODEL_SERVICE_URL
61           valueFrom:
62             configMapKeyRef:
63               name: my-config      # See ConfigMap defined at bottom
64               key: model.host
65         imagePullSecrets:
66         - name: ghcr-login-secret
```

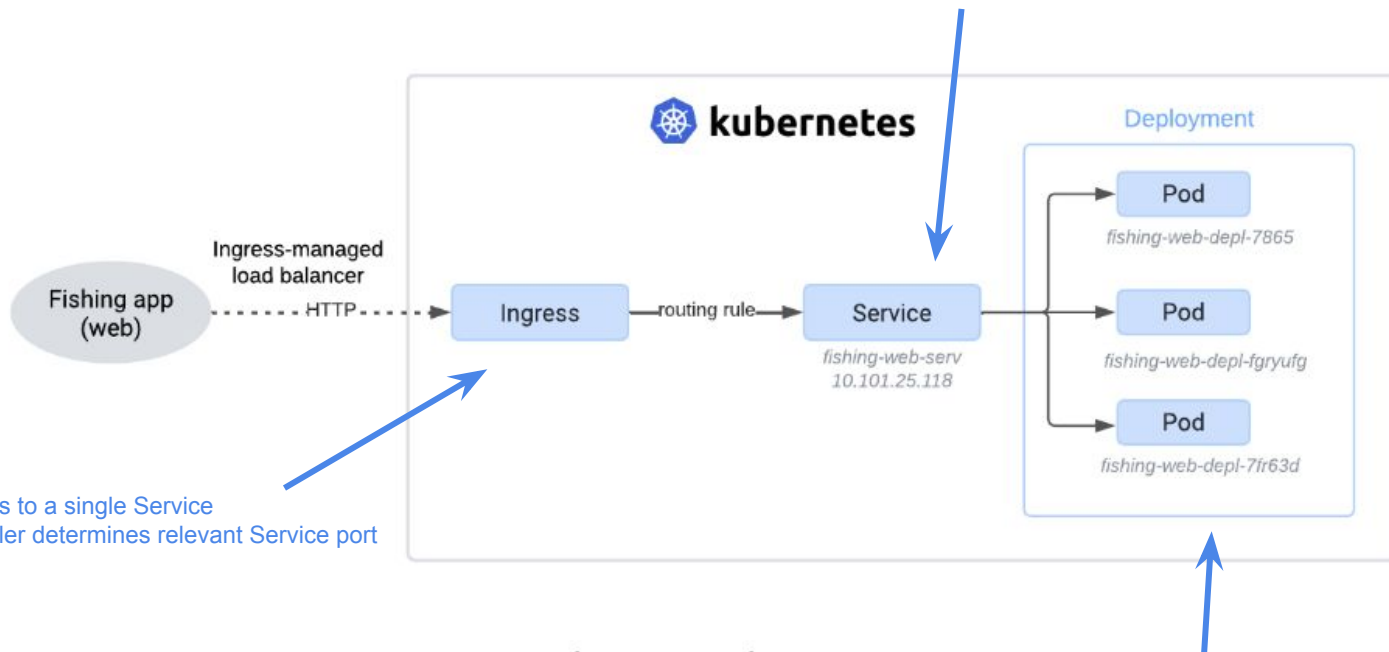


review

- Specification of desired state of Pods
- K8s ensures this state is retained + allows for scaling up or down replicas of Pods

4. K8s Deployment

- Prevent internal Pods from being accessed publicly
- Listen on specified port (fixed)
- Forward requests to containers in Pods



- Routes requests to a single Service
- Ingress Controller determines relevant Service port

Figure 2: Kubernetes Deployment Overview

- Specification of desired state of Pods
- K8s ensures this state is retained + allows for scaling up or down replicas of Pods

5. Experimental Setup

Requirements:

- Istio Service Mesh that deploys two versions of app-frontend and app-service.
 - Prometheus annotations to scrape.
- Verifiable hypothesis with a reject/validate criterion.
- Prometheus scrapes the relevant metrics.
 - Summary, counter, gauge, histogram.
 - Metrics about inference time, user interaction metrics, etc...
- Grafana displays results
- AlertManager sends an email for extreme cases (example: too many calls)

Hypothesis: “Given a version of the web-app that returns binary predictions (valid/phishing) to the user and a version that returns probabilities instead, the user will prefer the version of the web-app that returns binary predictions”

URL Phishing Detection

Library Version: v5.0.1

App version: v1

Enter Text:

google.com

Choose a model:

model.h5

Submit

This link is valid

Is this prediction accurate?

Yes

No

URL Phishing Detection

Library Version: v5.0.1

App version: v2

Enter Text:

google.com

Choose a model:

model.h5

Submit

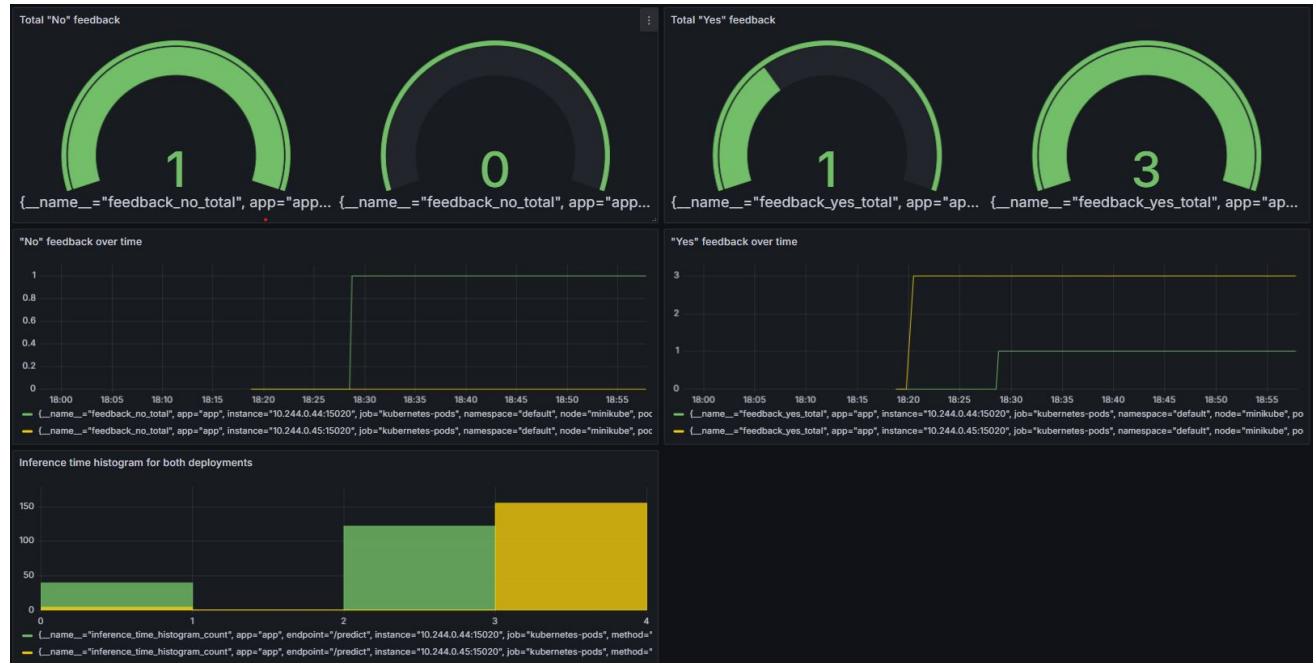
0.5110589

Is this prediction accurate?

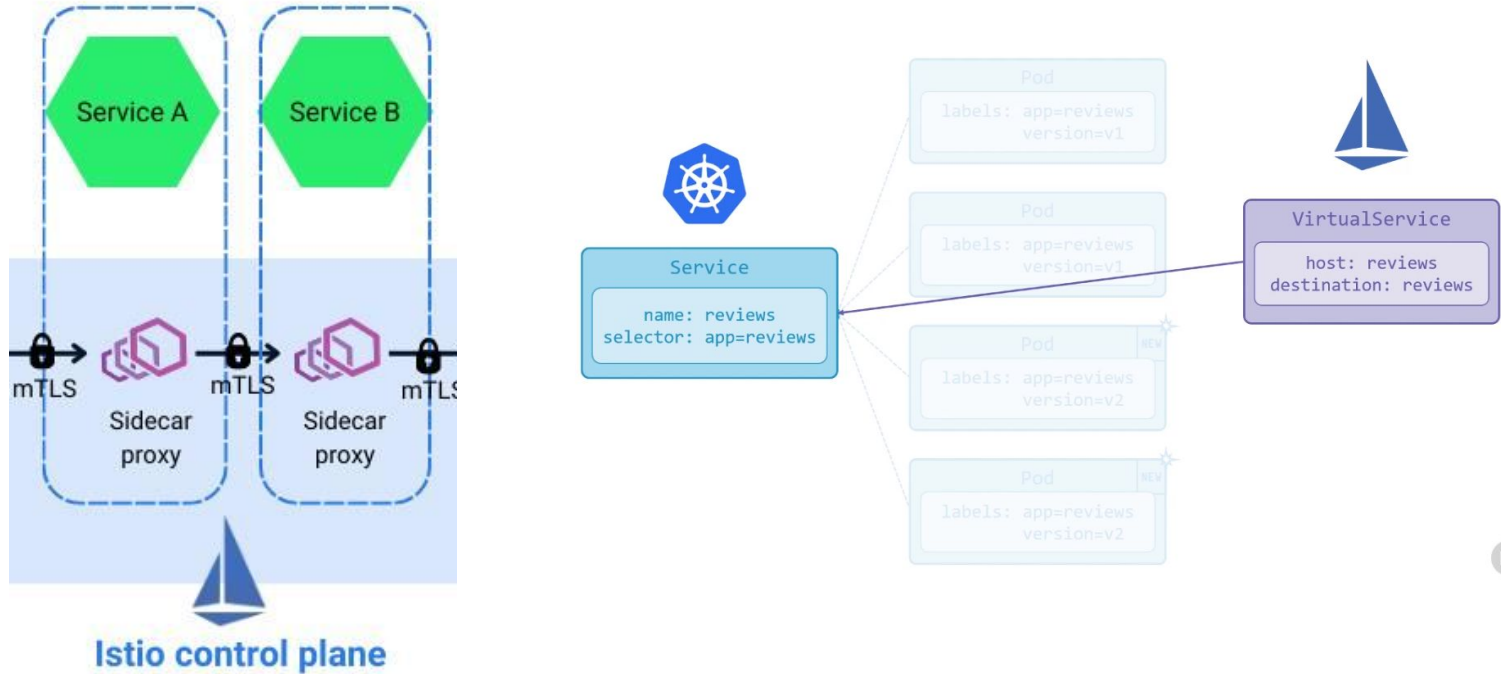
Yes

No

5. Experimental Setup: Prometheus and Grafana



5. Experimental Setup: Istio



6. Additional Istio Use Case-Rate Limiting

Egress Extension

In Istio, the Egress extension is used to manage and control traffic to external services in the service grid.

Adjustments to the Base Deployment

- Increase ENVOYFILTER configuration
- Enable and verify the flow limit function
- Monitoring and log analysis
- Adjustment and optimization strategy

```
PS C:\Users\maat\Desktop\REMLA-A5\ratelimit> kubectl exec -it sleep-0 -n rl -- sh
~ $ for i in $(seq 1 15); do curl -s -o /dev/null -w "%{http_code}\n" http://httpbin.r1:80/get; done
200
200
200
200
200
200
200
200
200
200
200
200
200
429
429
429
429
```

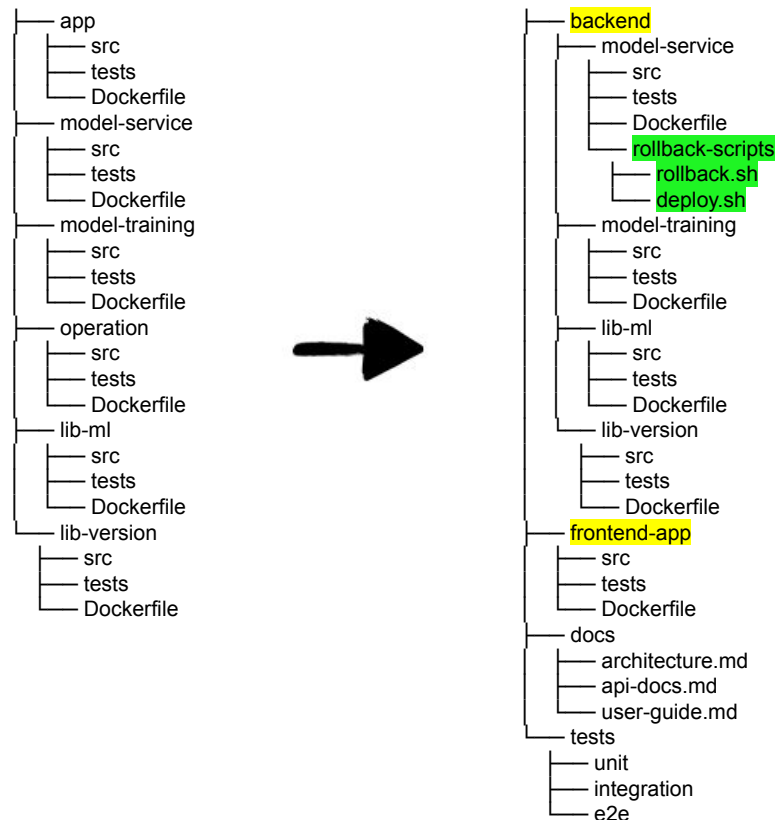
7. Extension Proposal- Rollback Strategy

Benefits:

- Improve System Stability
- Reduce deployment risk

Suggestions:

- Configure Automated Rollback with CI/CD Tools
- Blue-Green Deployment



Q&A

Thanks for listening! Questions?